

**Distributional or discrete? Social cues modulate the representations
underlying cross-situational learning**

Kyle MacDonald, Daniel Yurvosky, and Michael C. Frank

Department of Psychology, Stanford University

Author Note

We are grateful to the members of the Language and Cognition Lab for their feedback on this project. This work was supported by a National Science Foundation Graduate Research Fellowship to KM and an NIH NRSA Postdoctoral Fellowship to DY. Please address correspondence to Kyle MacDonald, Jordan Hall (Building 420), Stanford University, 450 Serra Mall, Stanford, CA 94305. Email: kyle.macdonald@stanford.edu.

Abstract

Word learning requires making inferences from noisy data – even concrete nouns occur in contexts with many possible referents. A statistical learner can handle this ambiguity by aggregating across naming events to form stable word-object mappings. But each naming event occurs within a social context that can vary along a continuum from low to high referential uncertainty. How does statistical word learning operate over observations that provide more or less evidentiary value? Drawing on social-pragmatic theories of language acquisition, we hypothesize that the presence of in-the-moment referential cues, like gaze, reduces referential uncertainty and modulates the underlying representation used for statistical learning. In three large-scale experiments with adults, we test the effect of varying referential uncertainty on attention and memory during cross-situational word learning. Referential cues shift learners away from tracking multiple hypotheses towards storing only a single hypothesis (Experiments 1 and 2). In addition, learners are sensitive to graded changes in the strength of a referential cue, and when it becomes less reliable, learners are more likely to store multiple hypotheses (Experiment 3). Together, the data suggest that representations in cross-situational word learning are quite flexible: In conditions of greater uncertainty, learners tend to store a broader range of information.

Keywords: statistical learning, pragmatic cues, word learning, language acquisition

Distributional or discrete? Social cues modulate the representations underlying cross-situational learning

Introduction

Learning the meaning of a new word should be hard. Consider that even concrete nouns are often used in complex contexts with multiple possible referents, which in turn have many conceptually natural properties that a speaker could talk about. This ambiguity creates the potential for an (in principle) unlimited amount of referential uncertainty in the learning task (Quine, 1960).¹ Remarkably, word learning proceeds despite this uncertainty, with estimates of adult vocabularies ranging between 50,000 to 100,000 distinct words (P. Bloom, 2002). How do learners infer and retain such a large variety of word meanings from data with this kind of ambiguity?

Statistical learning theories offer a solution to this learning problem by aggregating cross-situational statistics across labeling events to identify underlying word meanings (Siskind, 1996; Yu & Smith, 2007). Recent experimental work shows that both adults and young infants can use word-object co-occurrence statistics to learn words from individually ambiguous naming events (L. Smith & Yu, 2008; Vouloumanos, 2008). For example, L. Smith & Yu (2008) taught 12-month-olds three novel words simply by repeating consistent novel word-object pairings across 10 ambiguous exposure trials. Moreover, computational models suggest that cross-situational learning can scale up to learn adult-sized lexicons, even under conditions of considerable referential uncertainty (K. Smith, Smith, & Blythe, 2011).

¹This is a simplified version of Quine's *indeterminacy of reference* problem: That there are many possible meanings for a word ("Gavagai") that include the referent ("Rabbit") in their extension, e.g., "white," "rabbit," "dinner." Quine's broader philosophical point was that different meanings ("rabbit" and "undetached rabbit parts") could actually be extensionally identical and thus impossible to tease apart.

Although all cross-situational learning models agree that the input is the co-occurrence between words and objects and the output is stable word-object mappings, they disagree about how closely learners approximate the input distribution (for a review, see Smith, Suanda, & Yu (2014)). Some theories hold that we accumulate graded, statistical evidence about multiple referents for each word (McMurray, Horst, & Samuelson, 2012), while others argue that we track only a single candidate referent (Trueswell, Medina, Hafri, & Gleitman, 2013). Recent experimental and modeling work by Yurovsky & Frank (in press) suggests an integrative explanation: Learners allocate a fixed amount of their attention to one hypothesis, and the rest gets distributed evenly among the remaining alternatives. As the set of alternatives grows, the amount allocated to each object approaches zero.

Another ongoing debate in the literature is how to best characterize the input to cross-situational learning mechanisms. One way researchers have quantified ambiguity is to ask adults to guess the meaning of an intended referent from clips of caregiver-child interactions (Human Simulation Paradigm: HSP). Using the HSP, Medina, Snedeker, Trueswell, & Gleitman (2011) found that adults did not aggregate multiple word-referent correspondences across trials, concluding that real world learning contexts are too noisy to support tracking of multiple word-object links. In contrast, Yurovsky, Smith, & Yu (2013) found a bimodal distribution, with half of the naming episodes being unambiguous to adults and half being quite clear. In addition, Cartmill et al. (2013) showed that the proportion of unambiguous naming episodes varies across parents, with some parents' rarely providing highly informative contexts and others' doing so relatively often.

Thus, representations in cross-situational word learning can appear distributional or discrete, and the input to statistical learning mechanisms can vary along a continuum

from low to high ambiguity. These results raise an interesting question: could learners be sensitive to the ambiguity of the input and use this information to flexibly alter the representations they store in memory? In the current line of work, we investigate how the presence of pragmatic cues in the social context might alter the ambiguity of the input to statistical word learning mechanisms.

Social-pragmatic theories of language acquisition emphasize the importance of pragmatic cues for word learning (P. Bloom, 2002; Clark, 2009; Hollich et al., 2000). Experimental work shows that even children as young as 16 months are sophisticated intention-readers, preferring to map novel words to objects that are the target of a speaker’s gaze and not their own (Baldwin, 1993). In naturalistic observations, learners tend to retain labels that are accompanied with clear referential cues that are concurrent with visual access (Yu & Smith, 2012). And correlational data show strong links between early intention-reading skills (e.g., gaze following) and later vocabulary growth (Brooks & Meltzoff, 2005, 2008; Carpenter, Nagell, Tomasello, Butterworth, & Moore, 1998). Moreover, research outside the domain of language acquisition shows that the presence of social cues: (a) produces better spatial learning of audiovisual events (Wu, Gopnik, Richardson, & Kirkham, 2011), (b) boosts recognition of a cued object (Cleveland, Schug, & Striano, 2007), and (c) leads to preferential encoding of an object’s featural information (Yoon, Johnson, & Csibra, 2008). Together, the evidence suggests that social cues could help learners by allowing for efficient allocation of attention to the relevant statistics in the input, and thus change the representations stored in memory.

In the studies reported here, we ask whether the presence of a valid pragmatic cue, a speaker’s gaze, changes the representations underlying cross-situational word learning. We use a modified version of Yurovsky & Frank (in press)’s paradigm, which provides a direct measure of memory for alternative word-object links during cross-situational

learning. In Experiment 1, we manipulate the presence of a referential cue at different levels of attention and memory demands. At all levels of difficulty, learners tracked a strong single hypothesis, but learners were less likely to track multiple word-object links when a referential cue was present. In Experiment 2, we replicate the findings from Experiment 1 with a more ecologically valid social cue. In Experiment 3, we show that learners are sensitive to graded changes in the reliability of a referential cue and will flexibly increase the number of word-object links they store in response to changes in the quality of the input. In sum, the data suggest that cross-situational word learners are quite flexible, storing representations with different levels of fidelity depending on the amount of ambiguity present during learning.

Experiment 1

We set out to test the effects of a referential cue on cross-situational learning at different levels of attention and memory demands. Participants saw a series of ambiguous exposure trials that consisted of a set of novel objects (either 2, 4, 6, or 8) and an image of a schematic, female interlocutor. On each trial they heard one novel word that was either paired with a gaze cue or not, and were asked to make guesses about which object went with each word. In subsequent test trials, participants heard the novel word again after different numbers of intervening trials (0, 1, 3, and 7), this time paired with a new set of novel objects. Importantly, test trials were contingent upon participants' selection during exposure such that one of the objects in the set was either the participant's initial guess (Same trials) or one of the objects that was *not* the initial guess (Switch trials). While both single and multiple referent trackers could succeed on Same trials, only participants who encoded multiple word-object links during their first encounter could succeed on Switch trials. This provides a direct measure of

whether learners track multiple alternatives and if these representations are influenced by the presence of referential cues.

Method

Participants. This experiment was posted to Amazon Mechanical Turk as a set of Human Intelligence Tasks (HITs) to be completed only by participants with US IP addresses and an approval rate above 95%. Each HIT paid 30 cents. Approximately 50-130 HITs were posted for each of the 32 conditions (4 Referents X 4 Intervals X 2 Gaze conditions) for total of approximately 2400 paid HITs. If a participant completed the experiment more than once, he or she was paid each time but only data from the first HITs completion was included in the final data set. In addition, data was excluded from the final sample if participants did not give correct answers for familiar trials (5 HITs excluded).

Stimuli. Figure 1 shows stimuli used in Experiment 1. These stimuli consisted of black and white pictures of familiar and novel objects drawn from the set of 140 first used in Kanwisher, Woods, Iacoboni, & Mazziotta (1997), a schematic drawing of a human interlocutor, and audio recordings of familiar and novel words. Familiar words consisted of the labels for the familiar objects as produced by AT&T Natural Voices TM(voice: Crystal). Novel words were 1-3 syllable pseudowords obeying the rules of English phonotactics produced using the same speech synthesizer. A schematic drawing of a human speaker was chosen for ease of manipulating the direction of gaze, the referential cue of interest in this study. A sample experiment (as well as Experiments 2 and 3) can be viewed directly at the project page for this paper:
https://kemaconnald.github.io/soc_xsit/.

Design and Procedure. Participants were exposed to a series of 16 trials (8 exposure, and 8 test) in which they heard a speaker say one novel word, saw a set of novel objects, and were asked to guess which object went with the word. After a written explanation of the task, participants completed four practice trials that consisted of familiar words and objects. These trials also served to screen for participants who did not have their audio enabled or who were not attending to the task.

After the practice trials, participants were informed that they would now hear novel words, and see novel objects, and that they should continue selecting the correct referent for each word. Participants heard eight novel words twice, one exposure and one test trial for each word. Four of the test trials were Same trials in which the object that participants selected on the exposure trial appeared again amongst a set of new objects. The other four were Switch trials in which one of the objects in the set was selected randomly from the objects that the participant did not select on the previous exposure trial. All other objects were completely novel on each trial.

Participants were randomly assigned to see either 2, 4, 6, or 8 referents on each trial and to have either 0, 1, 3, or 7 trials in between exposure and test. Participants were also assigned to either the Gaze or No-gaze condition. In the Gaze condition, gaze was directed towards one of the objects on exposure trials; in the No-gaze condition, gaze was always directed straight ahead. On test trials, gaze was never informative. To indicate that participants' selections had been registered, a red dashed box appeared around the object they selected for one second after their click was received. This box appeared around the selected object whether or not it was the "correct" referent.

Results and Discussion

Analysis plan. The structure of our analysis plan is parallel across all three experiments. First, we examine performance on Exposure trials to provide evidence that learners were (a) sensitive to our manipulation and (b) altered their allocation of attention in response to changes in contextual ambiguity. Then we examine performance on Test trials to show that learners’ memory for alternative word-object links changes depending on the ambiguity of the learning context. The key behavioral prediction of our hypothesis is that the presence of gaze will result in reduced memory for multiple word-object links, operationalized as a decrease in performance on Switch test trials after seeing Exposure trials with a gaze cue. To quantify participants’ behavior, we use mixed effects regression models with the maximal random effects structure justified by our experimental design: by-subject intercepts and slopes for each trial type. All mixed-effects models were fit using the lme4 package in R (Bates, Maechler, Bolker, & Walker, 2013), and all of our data, processing, and analysis code can be viewed in the version control repository for this paper at:
https://github.com/kemacdonald/soc_xsit.

Exposure trials. To ensure that our referential cue manipulation was effective we compared participant’s performance on Exposure trials in the Gaze condition against the distribution expected if participants were selecting randomly (defined by a Binomial guessing model with four trials and a probability of success of $\frac{1}{NumReferents}$). Across all conditions, participants’ responses differed from those expected if participants were selecting randomly, exact binomial p (two-tailed) < .001, suggesting that gaze effectively directed participants’ attention to the target referent.

We were also interested in differences in participants’ response times across the experimental conditions. Since these trials were self-paced, participants could choose how much time to spend studying the referents on the screen, thus providing an index of

participants' attention. To quantify the effects of gaze, interval, and number of referents, we fit a linear mixed effects model predicting participants' response times as follows:

$$RT \sim \text{Gaze Condition} \times \text{Log(Interval)} \times \text{Log(Referents)} + (1 \mid \text{subject})$$

We found a significant main effect of referents ($\beta = 806.95$, $p < .001$) with slower responses as the number of referents increased, and a significant two-way interaction between Gaze condition and number of referents ($\beta = -517.43$, $p < .001$) such that responses were faster in the Gaze condition, especially as the number of referents increased. The interaction between Gaze condition and number of referents is shown in Panel A of Figure 2. Faster response times on Exposure trials with gaze provides preliminary evidence that the presence of a referential cue focused participants' attention on the gaze target and away from alternative word-object links.

Test trials. To analyze performance on test trials, we first compared the distribution of correct responses made by each participant to the distribution expected if participants were selecting randomly. Panel B of Figure 2 shows participants' accuracies in identifying the referent of each word in all conditions for both kinds of trials (Same and Switch) and in each referential cue condition (Gaze and No-gaze). We replicate the finding from Yurovsky & Frank (in press): at all referent and interval levels, both for Same and for Switch trials, participants' responses differed from those expected by chance (smallest $\chi^2_{2(4)} = 12.07$, $p < .01$). Participants' success on Switch trials provides direct evidence that learners encoded more than a single hypothesis in ambiguous word learning situations, even under high attentional and memory demands, and even in the presence of a valid referential cue.

To quantify the effect of each predictor on the probability of a correct response, we fit the following mixed-effects logistic regression model to a filtered dataset, removing

participants who were not reliably selecting the referent that was the target of gaze on exposure trials:²

$$Accuracy \sim Trial\ type \times Gaze \times Log(Interval) \times Log(Referents) + (Trial\ type \mid subject)$$

We follow Yurovsky & Frank (in press)’s analysis plan and coded interval and number of referents as continuous predictors and transformed these variables to the log scale. We limited the model to include only two-way interactions because the critical test of our hypothesis is the interaction between Gaze condition and Trial Type, and we did not have any theoretical predictions for possible three-way interactions.³

Table 1 shows the output of the logistic regression model. We found significant main effects of Referents ($\beta = -0.68$, $p < .001$) and Interval ($\beta = -0.59$, $p < .001$), such that as each of these factors increased, accuracy on test trials decreased. We also found significant main effects of Trial Type ($\beta = -1.58$, $p < .001$), with worse overall performance on Switch trials and in the Gaze condition ($\beta = -0.47$, $p < .05$). There were significant interactions between Trial Type and Interval ($\beta = 0.52$, $p < .001$), Trial Type and Referents ($\beta = -0.62$, $p < .001$), and Gaze condition and Referents ($\beta = 0.16$, $p < .05$). These interactions can be interpreted as (a) the interval between exposure and test affecting Same trials more than Switch trials, (b) the number of referents affecting

²We did not predict that there would be a subset of participants who would not follow the gaze cue, thus this filtering criteria was developed post-hoc. However, we believe the filter is theoretically motivated because we would only expect to see an effect of gaze if participants were actually using the gaze cue. The filter removes 90 participants who did not reliably select the gaze target on exposure trials. Importantly, the key inferences from the data do not depend on this filtering criteria.

³If we allow for three-way interactions, there is a significant interaction between Gaze condition, Trial Type, and Interval ($\beta = 0.25$, $p < .01$). Importantly, the two-way interaction between Gaze condition and Trial Type remains significant in this more complex model. A model including four-way interactions did not sufficiently improve model fit in order to justify the added complexity.

Switch trials more than Same trials, and (c) participants performing slightly better at higher number of referents in the Gaze condition (see Panel B of Figure 2). The interactions between Gaze condition and Referents and between Referents and Interval were not significant. Crucially, we found the predicted interaction between Trial Type and Gaze condition ($\beta = -0.55$, $p < .001$), with participants in the Gaze condition performing worse on Switch trials. This interaction provides direct evidence that the presence of a referential cue selectively reduced participants' memory for alternative word-object links.

Taken together, the response time and accuracy analyses provide evidence that the presence of a referential cue modulated learners' attention during learning, and in turn made them less likely to track multiple word-object links. Interestingly, we did not see strong evidence that reduced tracking of alternatives resulted in an increase in performance on Same trials. This finding suggests that the limitations on Same trials may be different than those regulating the distribution of attention on Switch trials, since the presence of a referential cue selectively reduced learners tracking of alternatives but apparently did not lead learners to form a stronger memory of their single candidate hypothesis.

It is important to point out that there was relatively large variation in performance across conditions both in group-level accuracy scores and in participants' tendency to *use* the referential cue on exposure trials. Moreover, we found a subset of participants who did not reliably use the gaze cue at all, potentially reducing the effect of gaze on cross-situational learning in this experiment. Perhaps, the effect of gaze was reduced because the referential cue that we used – a static schematic drawing of a speaker – was relatively weak compared to the cues present in real world learning environments. Hence, we do not yet know how learners' memory for alternatives during

cross-situational learning would change in the presence of a stronger and more ecologically valid referential cue. Experiment 2 attempts to answer this question.

Experiment 2

In Experiment 2, we attempt to replicate the findings from Experiment 1 using a more ecologically valid stimulus set. To move closer to a real world referential cue, we replaced the static, schematic drawing with a live actress. To reduce the across-conditions variability, we introduced a within-subjects design where each participant saw both Gaze and No-gaze exposure trials. We selected a subset of conditions from Experiment 1, testing only the 4-referent display with 0 and 3 intervening trials as between-subjects manipulations. Our goals were to replicate the reduction in learners' multiple alternatives tracking in the presence of referential cues, and to test whether increasing the ecological validity of the cue would result in a boost to the strength of learners' recall of their single candidate hypothesis.

Method

Participants. Participant recruitment and inclusionary/exclusionary criteria were identical to those of Experiment 1 (excluded 36 HITs). 100 HITs were posted for each condition (1 Referent X 2 Intervals X 2 Gaze conditions) for total of 400 paid HITs.

Stimuli. Audio and picture stimuli were identical to Experiment 1. The referential cue in the Gaze condition was a film of a live actress (see Figure 1). On each exposure trial, the actress looked out at the participant with a neutral expression, smiled, and then turned to look at one of the four images on the screen. She maintained her gaze for 3 seconds before returning to the center. On test trials, she looked straight ahead for the duration of the trial.

Design and Procedure

Procedures were identical to those of Experiment 1. The major design change was a within-subjects manipulation of the gaze cue with each participant seeing exposure trials with and without gaze. The experiment consisted of 32 trials broken down into 2 blocks of 16 trials. Each block consisted of 8 exposure trials and 8 test trials (4 Same trials and 4 Switch trials), and contained only Gaze or No-gaze exposure trials. The order of block was counterbalanced across participants.

Results and Discussion

We followed the same analysis plan as in Experiment 1, first analyzing performance on exposure trials, and then analyzing performance on test trials.

Exposure trials. Similar to Experiment 1, participants' responses on exposure trials differed from those expected by chance, exact binomial $p(\text{two-tailed}) < .001$, suggesting that gaze was effective in directing attention to the target referent. Participants in Experiment 2 were numerically more consistent in their use of gaze with the live action stimuli compared to the schematic stimuli used in Experiment 1 ($M1 = .76$, $M2 = .81$), suggesting that using a live actress resulted in a slight increase in participants' willingness to follow the gaze cue.

Panel A of figure 3 shows participants' response times. We replicate the findings from Experiment 1, with faster response times in Gaze condition. We fit a linear mixed effects model to response times with the same specification as Experiment 1, finding main effects for Gaze condition ($\beta = -1112.83$, $p < .001$) and Interval ($\beta = -498.96$, $p < .001$) with faster responses in the Gaze condition and in the longer Interval conditions. The two-way interaction between Gaze condition and interval was not significant, with gaze having the same effect on participants' response times at both intervals.

Test trials. Panel B of Figure 3 shows performance on test trials in Experiment 2. We replicate the critical finding from Experiment 1: after seeing exposure trials with gaze, participants performed worse on Switch trials, providing evidence that they stored fewer word-object links. We fit a mixed-effects logistic regression model with the same specifications as in Experiment 1 and found significant main effects of Interval ($\beta = -0.88$, $p < .001$) and Trial Type ($\beta = -2.74$, $p < .001$). Participants were less accurate as the interval between exposure and test increased and on the Switch trials overall.

In addition, there was a significant two-way interaction between Trial Type and Interval ($\beta = 0.76$, $p < .001$), with worse performance on Switch trials at the higher intervals, and a marginal two-way interaction between Gaze condition and Interval ($\beta = 0.13$, $p = 0.07$) such that the number of intervening trials had a smaller effect on participants' performance in the Gaze condition. Importantly, we found a robust interaction between Gaze condition and Trial Type ($\beta = -0.73$, $p < .001$) with Switch trials being more difficult after gaze exposure trials.⁴ In addition, we did not see evidence of a boost to performance on Same trials in the Gaze condition.

The results of Experiment 2 provide converging evidence for our hypothesis, showing that the presence of a referential cue reliably focused learners' attention away from alternative word-object links and shifted them towards single hypothesis tracking. Changing to a live action stimulus set led to slightly higher rates of selecting the target of gaze on exposure trials, but did not result in a boost to performance on Same trials.

⁴As in Experiment 1, we fit this model a filtered dataset removing participants who did not reliably use the gaze cue. We also fit the model to the unfiltered dataset, and we found no difference between the two analyses. This suggests that the combination of using stronger referential cue and switching to a within-subjects design reduced noise in our measurements.

The selective effect of gaze on Switch trials provides additional evidence that the fidelity of participants' single hypothesis was unaffected by the presence of a referential cue in our paradigm.

Thus far we have shown that people store different amounts of information in response to a categorical manipulation of referential uncertainty. In both Experiments 1 and 2, the learning context was either entirely ambiguous (No-gaze) or entirely unambiguous (Gaze). However, not all real world learning contexts fall at the extremes of this continuum (although see Yurovsky et al. (2013)). Could learners be sensitive to more subtle changes in the quality of learning contexts? In our next experiment, we test a clear prediction of our account: whether learners tendency to store multiple word-object links responds to graded changes of referential uncertainty in the learning context.

Experiment 3

In Experiment 3, we explore whether learners will allocate attention and memory flexibly in response to *graded* changes in the referential uncertainty present during learning. To test this hypothesis, we move beyond a categorical manipulation of the presence/absence of gaze, and we parametically vary the strength of the referential cue. We manipulate cue strength by including a block of familiarization trials at the start of the experiment where we establish the reliability of the speaker's gaze. This design was inspired by a growing body of experimental work showing that even young children are sensitive to the prior reliability of speakers and will use this information when deciding whom to learn novel words from (Koenig, Clement, & Harris, 2004).

Method

Participants. Participant recruitment, and inclusionary/exclusionary criteria were identical to those of Experiment 1 and 2 (excluded 4 HITs). 100 HITs were posted for each reliability level (0%, 25%, 50%, 75%, and 100%) for total of 500 paid HITs.

Design and Procedure. Procedures were identical to those of Experiment 1 and 2. We modified our cross-situational learning paradigm to include a block of 16 familiarization trials (8 exposure trials and 8 test trials), which established the reliability of the speaker. To establish reliability, we varied the proportion of Same/Switch trials that occurred during this familiarization block. Recall that on Switch trials the gaze target does not show up at test, thus providing evidence that this speaker’s gaze might not be a reliable cue to reference. Gaze reliability was a between-subjects manipulation, with participants either seeing 0, 2, 4, 6, or 8 Switch trials. After the familiarization block, participants completed another block of 16 trials (8 exposure trials and 8 test trials). Importantly, since we were no longer testing the effect of the presence or absence of a referential cue, all exposure trials in Experiment 3 included gaze, but this cue was more or less reliable depending on which familiarization block participants saw. Finally, at the end of the task we asked participants to assess the reliability of the speaker on a continuous scale from “completely unreliable” to “completely reliable.”

Results and Discussion

Exposure trials. Similar to Experiments 1 and 2, participants reliably chose the referent that was the target of gaze at rates greater than those that would be predicted by a guessing model $p(\text{two-tailed}) < .001$. To quantify the effect of reliability condition and participants’ subjective reliability assessment, we fit a mixed effects logistic regression model predicting the probability of selecting the gaze target (labeled as

“Accuracy”) as follows:

$$Accuracy \sim Reliability\ condition \times Subjective\ reliability + (1 \mid subject)$$

We found significant main effects of both reliability condition ($\beta = 3.3$, $p < .05$) and subjective reliability ($\beta = 7$, $p < .001$) such that when the gaze cue was more reliable and when subjective reliability assessments were higher participants were more likely to use the gaze cue. The interaction between speaker reliability and subjective reliability assessments was marginally significant ($\beta = -4.33$, $p = 0.1$). This suggests that participants were sensitive to the reliability manipulation both in how often they used the gaze cue during the task and in how they rated the speaker at the very end.

Test trials. Our primary prediction in this experiment is an interaction between reliability and test trial type, with higher levels of reliability leading to less attention and memory allocated to alternative word-object links. To test this prediction, we performed three complementary analyses of test trial performance using three different predictors: reliability condition, participants’ use of gaze, and participants’ subjective reliability assessment.

Reliability condition analysis. Panel A of Figure 4 shows participants’ accuracy on both types of test trials as a function of the reliability manipulation. We fit a mixed-effects logistic regression model predicting accuracy using reliability condition as a predictor and found a significant main effect of trial type ($\beta = -1.85$, $p < .001$), with lower accuracy on Switch trials. Importantly, we found a significant interaction between reliability and trial type ($\beta = -0.9$, $p = 0.05$), providing evidence for our key prediction. However, the interaction between our reliability manipulation and test trial performance was not very strong, and similar to Experiment 1 there was considerable variation across conditions (see the 50% reliable condition in Panel A of Figure 4).

Thus, we conducted a follow-up analysis where we modeled accuracy on test trials as a function of how often participants used the gaze cue on exposure trials.

Gaze use analysis. We would only expect to see a strong interaction between reliability and trial type if learners chose to use the gaze cue during exposure trials. To test this hypothesis, we fit a mixed effects logistic regression model with the same specifications, but substituting accuracy on exposure trials for reliability condition as a predictor. We found a robust two-way interaction between accuracy on exposure trials and Trial Type ($\beta = -0.36$, $p < .001$) such that participants who were more likely to use the gaze cue performed worse on Switch trials, but not Same trials.⁵ Panel B of Figure 4 shows this interaction.

Subjective reliability analysis. The strong interaction between frequency of gaze use and test trial performance suggests that participants’ subjective experience of reliability in the experiment mattered. Thus, we fit the same mixed effects logistic regression model, but substituted subjective reliability as a predictor of test trial performance. We found a significant interaction between Trial Type and participants’ subjective reliability assessments ($\beta = -1.37$, $p = 0.02$) – if participants thought the speaker was more reliable, then they performed worse on Switch, but not Same, trials.

Taken together, these three analyses show that as the speaker’s gaze became more reliable, participants were (a) more likely to use it, (b) more likely to rate the speaker as reliable, and (c) less likely to store multiple word-object links. These findings support and extend the results of Experiments 1 and 2 in several important ways. First, participants’ performance on Same trials was again relatively unaffected by changes in

⁵We found this interaction while performing exploratory data analysis on a previous version of this study ($N = 250$, $\beta = -0.28$, $p < .001$). The results reported here are from a follow-up study where testing this interaction was a planned analysis.

performance on Switch trials. This provides converging evidence that the limitations on Same trials may be different than those regulating the distribution of attention on Switch trials. Second, learners' *use* of a referential cue was a stronger predictor of reduced memory for alternative word-object links compared to our reliability manipulation. It is important to note that although we found a significant effect of reliability on participants' use of the gaze cue, participants' tendency to use the cue remained high. Consider that even in the 0% reliability condition the mean proportion of gaze following was still 0.82. It is reasonable that participants would continue to use the gaze cue in our experiment since it was the only cue available and participants did not have a strong reason to think that the speaker would be deceptive.

The critical contribution of Experiment 3 is the finding that learners respond to a *graded* manipulation of referential uncertainty, with the amount of information stored from the initial exposure tracking with both the reliability of the cue and participants' use of the cue. This provides support for a critical prediction of our account: that learners store a strong single candidate word meaning and a set of alternatives with different levels of fidelity that corresponded with the amount of referential uncertainty present during the initial exposure to a word.

General Discussion

Tracking cross-situational word-object statistics allows word learning to proceed even in the presence of individually ambiguous naming events. But models of cross-situational learning disagree about how much information is actually stored in memory and how to best characterize the input to statistical learning mechanisms. In the current line of work, we explore the hypothesis that these two factors are fundamentally linked, both to one another and to the social context in which word learning occurs.

Specifically, we ask how cross-situational learning operates over social input that varies along a continuum from low to high ambiguity.

Our results suggest that the representations underlying cross-situational learning are quite flexible. In the absence of a referential cue to word meaning, learners were more likely to store alternative word-object links. In contrast, when gaze was present, they stored less information, showing behavior consistent with tracking a single hypothesis (Experiments 1 and 2). Learners were also sensitive to a parametric manipulation of the referential cue, showing a graded increase in the tendency to use the cue as reliability increased, which in turn resulted in a graded decrease in memory for alternative word-object links (Experiment 3). Interestingly, across all three experiments, reduced memory for alternative hypotheses did not result in a boost in memory for learners' candidate hypothesis. This pattern of data suggests that the presence of a referential cue selectively affected one component of the underlying representation: the number of alternative word-object links, and not learners candidate hypothesis.

Why did we not see an increase in the strength of learners' candidate hypothesis? One possibility is that participants did not shift their cognitive resources from the set of alternatives to their single hypothesis, but instead rationally conserved their resources for future use. Griffiths, Lieder, & Goodman (2015) formalize this behavior by pushing the rationality of computational-level models down to the psychological process level. In their framework, cognitive systems are thought to be adaptive in that they optimize the use of their limited resources, taking the cost of computation (e.g., opportunity cost of time or mental opportunity) into account. For example, Vul, Goodman, Griffiths, & Tenenbaum (2014) showed that as time pressure increased in a decision-making task, participants were more likely to show behavior consistent with a less cognitively challenging strategy of matching, rather than with the globally optimal strategy. Here,

we show evidence that learners adapt their allocation of cognitive resources to the level of referential uncertainty in the learning context, spending less time studying alternative word-object links and reducing the number of links stored in memory when uncertainty is low.

Our results also fit well with recent experimental work that investigates how attention and memory can constrain infants' statistical word learning. For example, Smith & Yu (2013) used a modified cross-situational learning task to show that only infants who disengaged from a novel object to look at both potential referents were able to learn the correct word-object mappings. Moreover, Vlach & Johnson (2013) showed that 16-month-olds were only able to learn from adjacent cross-situational co-occurrence statistics, and unable to learn from co-occurrences that were separated in time. Both of these findings make the important point that only the data that comes into contact with the learning system can be used for cross-situational word learning, and this data is directly influenced by the attention and memory constraints of the learner. Our findings suggest that referential cues could play an important role in constraining the input to statistical learning mechanisms.

How should we characterize the effect of social information on attention and memory in our task? One possibility is that the referential cue acts as a filter, only allowing likely referents to contact statistical learning mechanisms (Yu & Ballard, 2007). This 'filtering account' separates the effect of social cues from the underlying computation that aggregates cross-situational information. Another possibility is that referential cues provide evidence about a speaker's communicative intent (Frank, Goodman, & Tenenbaum, 2009). In this model, the learner is reasoning about the speaker and word meanings simultaneously, which places inferences based on social information as part of the underlying computation. A third possibility is that

participants thought of the referential cue as pedagogical. In this scenario, learners assume that the speaker will choose an action that is most likely to increase the learner's belief in the true state of the world (Shafto, Goodman, & Frank, 2012), making it unnecessary to allocate resources to alternative hypotheses. Experiments show that children spend less time exploring an object and are less likely to discover alternative object-functions, if a single function is demonstrated in a pedagogical context (Bonawitz et al., 2011). However, because the results from the current study cannot distinguish between these explanations, these questions remain topics for future studies specifically designed to tease apart these possibilities.

There are several limitations to the current study that are worth noting. First, the social context we used was relatively impoverished. Although we moved beyond a simple manipulation of the presence or absence of social information, we isolated just a single cue to reference, gaze. But real-world learning contexts are much more complex, providing learners access to multiple cues such as gaze, pointing, and previous discourse. In fact, Frank, Tenenbaum, & Fernald (2013) analyzed a corpus of parent-child interactions and concluded that learners would do better to aggregate noisy social information from multiple cues, rather than monitor a single cue, because no single cue was a consistent predictor of reference in their corpus. In our data, we did see a more reliable effect of referential cues when we used a live actress, which included both gaze and head turn as opposed to the static, schematic stimuli, which only included gaze. It is still an open and interesting question as to how our results would generalize to real-world learning environments that contain a rich combination of social cues.

Second, we do not yet know how these results would generalize to young word learners. Research with infants' shows rapid development of visual attention and memory in the first years of life (Colombo, 2001; Ross-sheehy, Oakes, & Luck, 2003).

Moreover, experimental work shows that infants' attention is often stimulus driven and sticky (Oakes, 2011), suggesting that very young word learners might not effectively explore the visual scene to extract the necessary statistics for effective cross-situational word learning. The current work suggests that referential cues might play an even more important role for young learners, guiding them to the relevant statistics in the input.

And third, in the current experiments we tested a minimal cross-situational learning scenario. Our task contained only one exposure for each novel word-object pairing. In contrast, real world naming events are best characterized by discourse, where an object is likely to be named repeatedly in a short amount of time (Frank et al., 2013; Rohde & Frank, 2014). Moreover, we presented novel words in isolation, removing any sentential cues to word meaning (e.g., verb-argument relations). Previous work shows that sentence-level constraints can interact with and complement cross-situational word learning mechanisms (Koehne & Crocker, 2014). Thus, we need more evidence to understand how representations underlying cross-situational learning change in response to referential uncertainty at different timescales and in richer language contexts that more accurately reflect learning environments.

Word learning proceeds despite the potential for high levels of referential uncertainty and learners' limited cognitive resources. Our work shows that cross-situational learners flexibly respond to the amount of ambiguity in the input, and as referential uncertainty increases, learners store more word-object links. Overall, these results bring together aspects of both social and statistical accounts of word learning, and increase our understanding of how statistical learning mechanisms operate over fundamentally social input.

References

- Baldwin, D. A. (1993). Infants' ability to consult the speaker for clues to word reference. *Journal of Child Language*, 20(02), 395–418.
- Bates, D., Maechler, M., Bolker, B., & Walker, S. (2013). Lme4: Linear mixed-effects models using eigen and s4. *R Package Version*, 1(4).
- Bloom, P. (2002). *How children learn the meaning of words*. The MIT Press.
- Bonawitz, E., Shafto, P., Gweon, H., Goodman, N. D., Spelke, E., & Schulz, L. (2011). The double-edged sword of pedagogy: Instruction limits spontaneous exploration and discovery. *Cognition*, 120(3), 322–330.
- Brooks, R., & Meltzoff, A. N. (2005). The development of gaze following and its relation to language. *Developmental Science*, 8(6), 535–543.
- Brooks, R., & Meltzoff, A. N. (2008). Infant gaze following and pointing predict accelerated vocabulary growth through two years of age: A longitudinal, growth curve modeling study. *Journal of Child Language*, 35(01), 207–220.
- Carpenter, M., Nagell, K., Tomasello, M., Butterworth, G., & Moore, C. (1998). Social cognition, joint attention, and communicative competence from 9 to 15 months of age. *Monographs of the Society for Research in Child Development*, i–174.
- Cartmill, E. A., Armstrong, B. F., Gleitman, L. R., Goldin-Meadow, S., Medina, T. N., & Trueswell, J. C. (2013). Quality of early parent input predicts child vocabulary 3 years later. *Proceedings of the National Academy of Sciences*, 110(28), 11278–11283.
- Clark, E. V. (2009). *First language acquisition*. Cambridge University Press.
- Cleveland, A., Schug, M., & Striano, T. (2007). Joint attention and object learning in 5-and 7-month-old infants. *Infant and Child Development*, 16(3), 295–306.
- Colombo, J. (2001). The development of visual attention in infancy. *Annual Review of Psychology*, 52(1), 337–367.

Frank, M. C., Goodman, N. D., & Tenenbaum, J. B. (2009). Using speakers' referential intentions to model early cross-situational word learning. *Psychological Science*, 20(5), 578–585.

Frank, M. C., Tenenbaum, J. B., & Fernald, A. (2013). Social and discourse contributions to the determination of reference in cross-situational word learning. *Language Learning and Development*, 9(1), 1–24.

Griffiths, T. L., Lieder, F., & Goodman, N. D. (2015). Rational use of cognitive resources: Levels of analysis between the computational and the algorithmic. *Topics in Cognitive Science*, 7(2), 217–229.

Hollich, G. J., Hirsh-Pasek, K., Golinkoff, R. M., Brand, R. J., Brown, E., Chung, H. L., . . . Bloom, L. (2000). Breaking the language barrier: An emergentist coalition model for the origins of word learning. *Monographs of the Society for Research in Child Development*, i–135.

Kanwisher, N., Woods, R. P., Iacoboni, M., & Mazziotta, J. C. (1997). A locus in human extrastriate cortex for visual shape analysis. *Journal of Cognitive Neuroscience*, 9(1), 133–142.

Koehne, J., & Crocker, M. W. (2014). The interplay of cross-situational word learning and sentence-level constraints. *Cognitive Science*.

Koenig, M. A., Clement, F., & Harris, P. L. (2004). Trust in testimony: Children's use of true and false statements. *Psychological Science*, 15(10), 694–698.

McMurray, B., Horst, J. S., & Samuelson, L. K. (2012). Word learning emerges from the interaction of online referent selection and slow associative learning. *Psychological Review*, 119(4), 831.

Medina, T. N., Snedeker, J., Trueswell, J. C., & Gleitman, L. R. (2011). How words can and cannot be learned by observation. *Proceedings of the National Academy*

of Sciences, 108(22), 9014–9019.

Oakes, L. M. (2011). *Infant perception and cognition: Recent advances, emerging theories, and future directions*. Oxford University Press, USA.

Quine, W. V. (1960). 0. word and object. *111e MIT Press*.

Rohde, H., & Frank, M. C. (2014). Markers of topical discourse in child-directed speech. *Cognitive Science*, 38(8), 1634–1661.

Ross-sheehy, S., Oakes, L. M., & Luck, S. J. (2003). The development of visual short-term memory capacity in infants. *Child Development*, 74(6), 1807–1822.

Shafto, P., Goodman, N. D., & Frank, M. C. (2012). Learning from others the consequences of psychological reasoning for human learning. *Perspectives on Psychological Science*, 7(4), 341–351.

Siskind, J. M. (1996). A computational study of cross-situational techniques for learning word-to-meaning mappings. *Cognition*, 61(1), 39–91.

Smith, K., Smith, A. D., & Blythe, R. A. (2011). Cross-situational learning: An experimental study of word-learning mechanisms. *Cognitive Science*, 35(3), 480–498.

Smith, L. B., & Yu, C. (2013). Visual attention is not enough: Individual differences in statistical word-referent learning in infants. *Language Learning and Development*, 9(1), 25–49.

Smith, L. B., Suanda, S. H., & Yu, C. (2014). The unrealized promise of infant statistical word-referent learning. *Trends in Cognitive Sciences*, 18(5), 251–258.

Smith, L., & Yu, C. (2008). Infants rapidly learn word-referent mappings via cross-situational statistics. *Cognition*, 106(3), 1558–1568.

Trueswell, J. C., Medina, T. N., Hafri, A., & Gleitman, L. R. (2013). Propose but verify: Fast mapping meets cross-situational word learning. *Cognitive Psychology*, 66(1), 126–156.

- Vlach, H. A., & Johnson, S. P. (2013). Memory constraints on infants' cross-situational statistical learning. *Cognition*, *127*(3), 375–382.
- Vouloumanos, A. (2008). Fine-grained sensitivity to statistical information in adult word learning. *Cognition*, *107*(2), 729–742.
- Vul, E., Goodman, N., Griffiths, T. L., & Tenenbaum, J. B. (2014). One and done? Optimal decisions from very few samples. *Cognitive Science*, *38*(4), 599–637.
- Wu, R., Gopnik, A., Richardson, D. C., & Kirkham, N. Z. (2011). Infants learn about objects from statistics and people. *Developmental Psychology*, *47*(5), 1220.
- Yoon, J. M., Johnson, M. H., & Csibra, G. (2008). Communication-induced memory biases in preverbal infants. *Proceedings of the National Academy of Sciences*, *105*(36), 13690–13695.
- Yu, C., & Ballard, D. H. (2007). A unified model of early word learning: Integrating statistical and social cues. *Neurocomputing*, *70*(13), 2149–2165.
- Yu, C., & Smith, L. B. (2007). Rapid word learning under uncertainty via cross-situational statistics. *Psychological Science*, *18*(5), 414–420.
- Yu, C., & Smith, L. B. (2012). Embodied attention and word learning by toddlers. *Cognition*.
- Yurovsky, D., & Frank, M. C. (in press). An integrative account of constraints on cross-situational learning. *Cognition*.
- Yurovsky, D., Smith, L. B., & Yu, C. (2013). Statistical word learning at scale: The baby's view is better. *Developmental Science*, *16*(6), 959–966.

Table 1

Predictor estimates with standard errors and significance information for a logistic mixed-effects model predicting word learning in Experiment 1.

Predictor	Estimate	Std. Error	<i>z</i> value	<i>p</i> value	
Intercept	4.16	0.20	20.39	< .001	***
Switch Trial	-1.58	0.17	-9.29	< .001	***
Gaze Condition	-0.47	0.21	-2.19	0.03	*
Log(Interval)	-0.59	0.08	-7.34	< .001	***
Log(Referents)	-0.68	0.08	-8.36	< .001	***
Switch Trial : Gaze Condition	-0.55	0.10	-5.45	< .001	***
Switch Trial : Log(Interval)	0.52	0.04	11.93	< .001	***
Switch Trial : Log(Referent)	-0.62	0.06	-9.67	< .001	***
Gaze Condition : Log(Interval)	0.08	0.05	1.53	0.13	
Gaze Condition : Log(Referent)	0.16	0.08	2.00	0.05	*
Log(Interval) : Log(Referent)	-0.00	0.03	-0.05	0.96	

Table 2

Predictor estimates with standard errors and significance information for a logistic mixed-effects model predicting word learning in Experiment 1.

Predictor	Estimate	Std. Error	<i>z</i> value	<i>p</i> value	
Intercept	2.69	0.16	16.84	< .001	***
Switch Trial	-2.74	0.16	-17.04	< .001	***
Gaze Condition	-0.12	0.16	-0.75	0.45	
Log(Interval)	-0.88	0.09	-9.40	< .001	***
Switch Trial : Gaze Condition	-0.73	0.15	-4.85	< .001	***
Switch Trial : Log(Interval)	0.76	0.09	8.35	< .001	***
Gaze Condition : Log(Interval)	0.13	0.07	1.81	0.07	.




Figure 1. Examples of exposure and test trials from Experiment 1 (schematic gaze cue) and Experiments 2 & 3 (human actress gaze cue). Participants saw exposure trials with or without a gaze cue depending on condition assignment. All participants saw both types of test trials: same and switch. On Same trials the object that participants chose during exposure appeared with a new novel object. On switch trials the object that participants did not choose appeared with a new novel object.

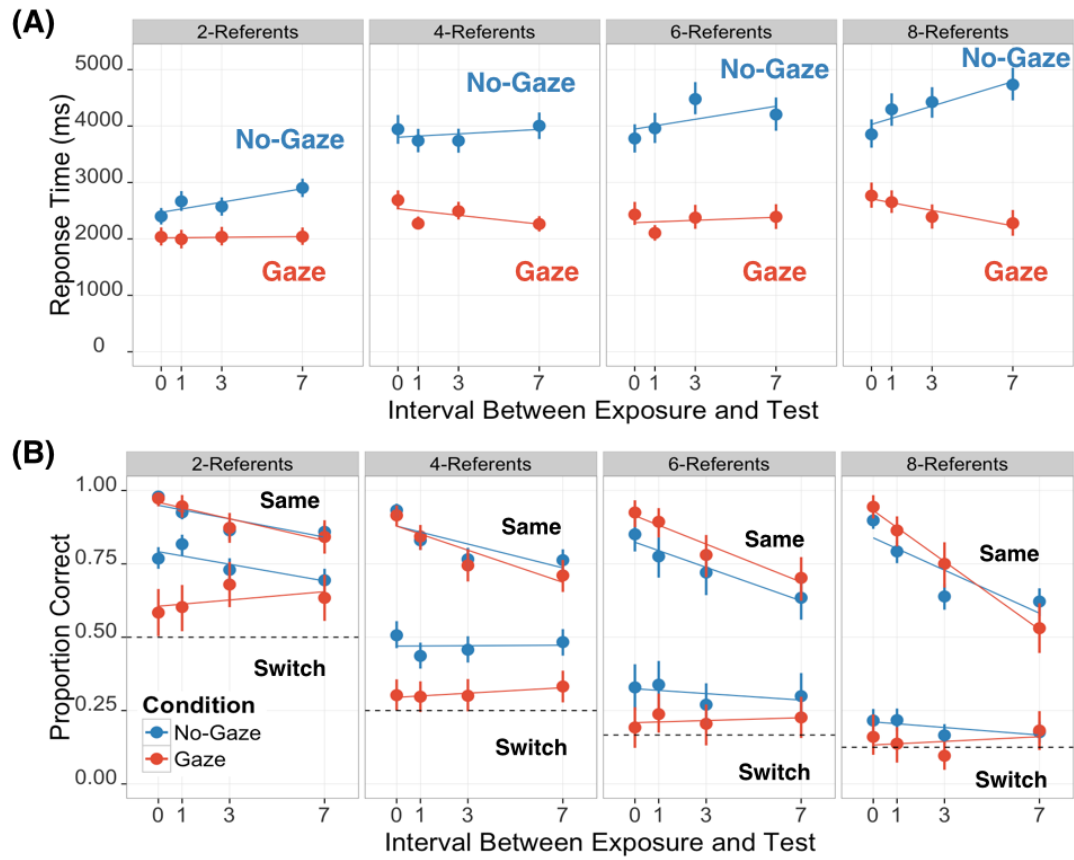


Figure 2. Experiment 1 results. Panel A shows response times on exposure trials across all experimental conditions: Gaze and No-gaze, Referents (2, 4, 6, and 8), and Intervening trials (0, 1, 3, and 7). Panel B shows accuracy on test trials for both trial types (Same and Switch) across all conditions. The horizontal dashed lines represent chance performance for each condition. Colored lines are linear model fits and error bars indicate 95% confidence intervals computed by non-parametric bootstrap.

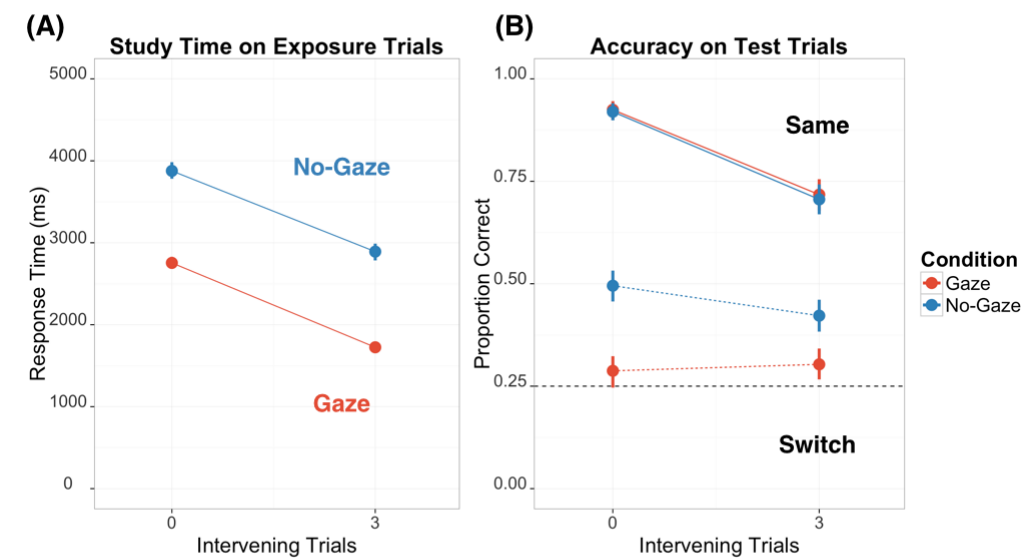


Figure 3. Experiment 2 results. Panel A shows study times for exposure trials with and without gaze. Panel B shows accuracy on test trials for same and Switch trials across all conditions. The dashed line in Panel B represents chance performance. Error bars indicate 95% confidence intervals computed by non-parametric bootstrap.

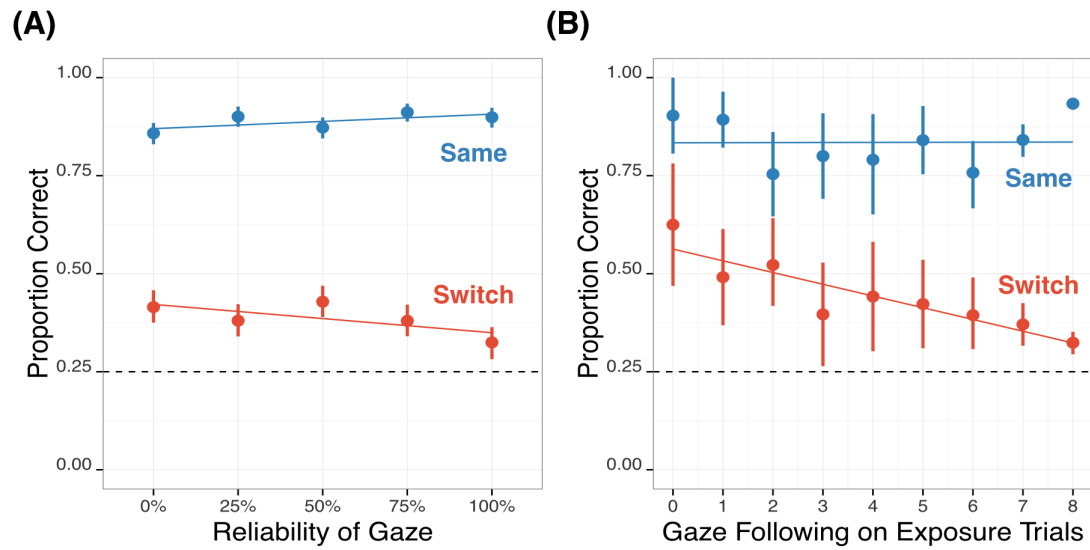


Figure 4. Accuracy on test trials in Experiment 3 for both same and switch trial types. Panel A shows accuracy as a function of the speaker’s reliability. Panel B shows accuracy as a function of participants’ gaze following on exposure trials. The horizontal dashed line represents the expected performance if participants were selecting randomly. The colored lines are linear model fits and error bars indicate 95% confidence intervals computed by non-parametric bootstrap.